

KMEANS

```
library(plyr)
library(ggplot2)
library(cluster)
library(lattice)
library(graphics)
library(grid)
library(gridExtra)



#import the student grades
grade_input=as.data.frame(read.csv("D:/DSA_DATASETS/datasets/grades_km_input.csv"))
kmdata_orig=as.matrix(grade_input[,c("Student","English","Math","Science")])

kmdata<-kmdata_orig[,2:4]
kmdata[1:10,]
wss<-numeric(15)

for(k in 1:15) wss[k]<-sum(kmeans(kmdata,centers=k,nstart=25) $ withinss)
plot(1:15,wss,type="b",xlab="Number of clusters",ylab="Within Sum of Squares")
km=kmeans(kmdata,3,nstart=25)
km
c(wss[3],sum(km$withinss))
df=as.data.frame(kmdata_orig[,2:4])
df$cluster=factor(km$cluster)
km$cluster
centers=as.data.frame(km$centers)
ggplot()
g1= ggplot(data=df,aes(x=English,y=Math,color=cluster)) +
```

```
geom_point() + theme(legend.position="right") +  
geom_point(data=centers,aes(x=English,y=Math,color=as.factor(c(1,2,3))),  
           size=10,alpha=.3,show.legend = FALSE)  
g2 =ggplot(data=df,aes(x=English,y=Science,color=cluster )) +  
geom_point() +  
geom_point(data=centers,aes(x=English,y=Science,color=as.factor(c(1,2,3))),  
           size=10,alpha=.3,show.legend=FALSE)  
g3 =ggplot(data=df,aes(x=Math,y=Science,color=cluster)) +  
geom_point() +  
geom_point(data=centers,aes(x=Math,y=Science, color=as.factor(c(1,2,3))),  
           size=10,alpha=.3,show.legend = FALSE)  
tmp = ggplot_gtable(ggplot_build(g1))  
library(grid)  
library(gridExtra)  
grid.arrange(g1,g2,g3,ncol=1,top="High School Student Cluster Analysis")
```

OUTPUT:



Plots

```
# Load required libraries
```

```
library(ggplot2)
```

```
library(GGally)
```

```
# Load dataset
```

```
data <- mtcars
```

```
# Scatter Plot
```

```
ggplot(mtcars, aes(x = hp, y = mpg)) + geom_point()
```

```
# Boxplot
```

```
ggplot(mtcars, aes(x = factor(cyl), y = mpg)) + geom_boxplot()
```

```
# Bar Chart
```

```
ggplot(mtcars, aes(x = factor(cyl))) + geom_bar()
```

```
# Line Chart
```

```
ggplot(mtcars, aes(x = hp, y = mpg)) + geom_line()
```

```
# Hexbin Plot (requires hexbin package)
```

```
if (!requireNamespace("hexbin", quietly = TRUE)) install.packages("hexbin")
```

```
ggplot(mtcars, aes(x = hp, y = mpg)) + geom_hex()
```

```
# Dot plot
```

```
ggplot(mtcars, aes(x = mpg)) + geom_dotplot(binwidth = 0.5)
```

```
# Histogram, Density, and Rug Plot
```

```
income <- rlnorm(4000, meanlog = 4, sdlog = 0.7)
```

```
summary(income)
```

```
# Convert income to thousands
```

```
income <- 1000 * income
```

```
summary(income)
```

```
# Histogram
```

```
hist(income, breaks = 500, xlab = "Income", main = "Histogram of Income")
```

```
# Density plot
```

```
plot(density(log10(income), adjust = 0.5),  
     main = "Distribution of Income (log10 scale)")
```

```
# Add rug to the density plot
```


```
rug(log10(income))
```

```
# Scatter Plot Matrix using GGally
```


```
ggpairs(data[, c("mpg", "hp", "wt")])
```



OUTPUTS:



Histogram of Income





Distribution of Income (log10 scale)

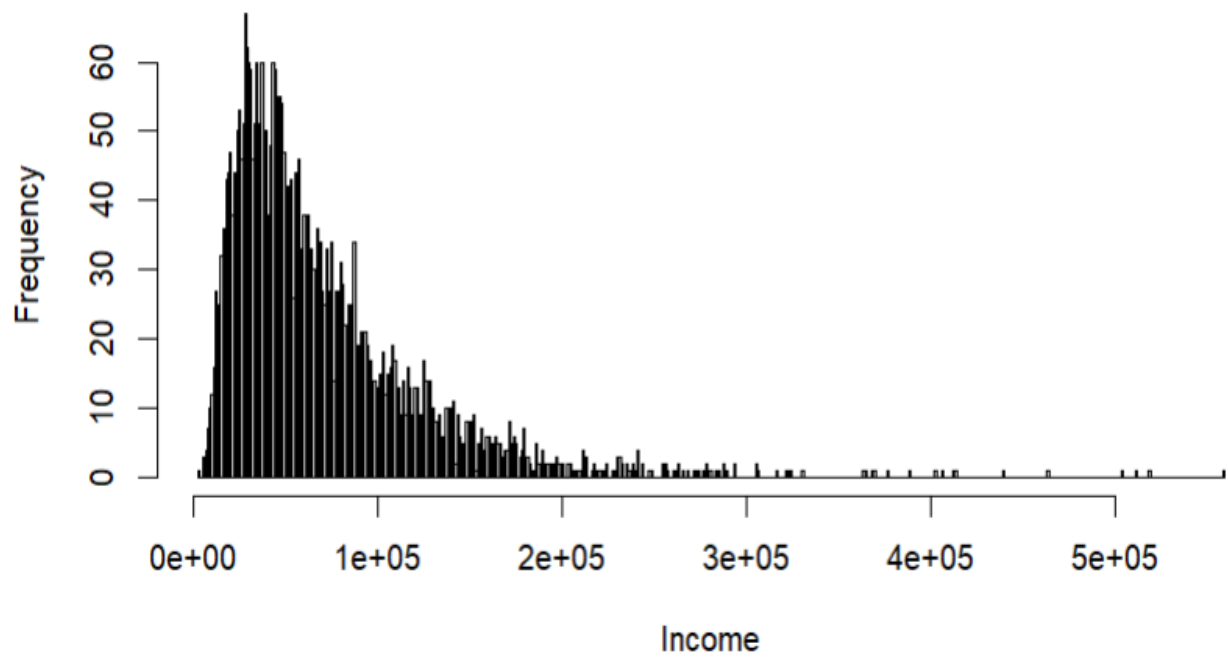







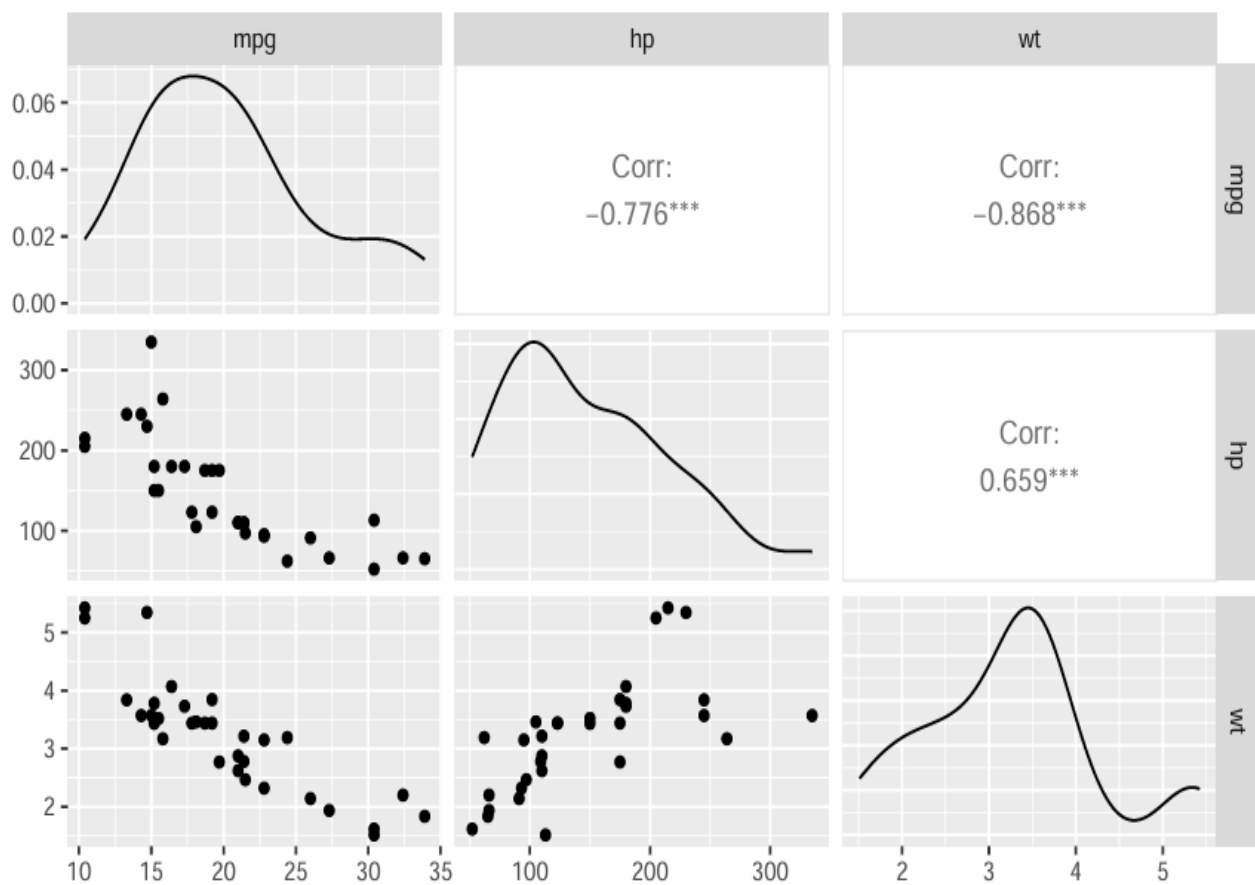


Histogram of Income



Distribution of Income (log10 scale)





LOGISTIC REGRESSION

```
#####
```

```
# This code covers the code presented in
```

```
# Section 6.2 Logistic Regression
```

```
#####
```

```
#####
```

```
# Section 6.2.3 Diagnostics
```

```
#####
```

```
churn_input = as.data.frame(read.csv("C:/Users/vanif/Desktop/churn.csv"))
```

```
head(churn_input) #displays first few rows n columns from the dataset
```

```
sum(churn_input$Churned)
```

```
Churn_logistic1 <- glm (Churned~Age + Married + Cust_years + Churned_contacts,  
                        data=churn_input, family=binomial(link="logit"))#"glm()=Generalized Linear  
Models"
```

```
summary(Churn_logistic1)
```

```
Churn_logistic2 <- glm (Churned~Age + Married + Churned_contacts,
```

```
                        data=churn_input, family=binomial(link="logit"))
```

```
summary(Churn_logistic2)
```

```
Churn_logistic3 <- glm (Churned~Age + Churned_contacts,
```

```
                        data=churn_input, family=binomial(link="logit"))
```

```
summary(Churn_logistic3)
```

```
# Deviance and the Log Likelihood Ratio Test
```

```

# Using the residual deviances from Churn_logistics2 and Churn_logistic3
# determine the significance of the computed test statistic
summary(Churn_logistic2)
pchisq(.9 , 1, lower=FALSE)

# Receiver Operating Characteristic (ROC) Curve

#install.packages("ROCR")    #install, if necessary
library(ROCR)

pred = predict(Churn_logistic3, type="response")
predObj = prediction(pred, churn_input$Churned )

rocObj = performance(predObj, measure="tpr", x.measure="fpr")
aucObj = performance(predObj, measure="auc")

plot(rocObj, main = paste("Area under the curve:", round(aucObj@y.values[[1]],4)))

# extract the alpha(threshold), FPR, and TPR values from rocObj
alpha <- round(as.numeric(unlist(rocObj@alpha.values)),4)
fpr <- round(as.numeric(unlist(rocObj@x.values)),4)
tpr <- round(as.numeric(unlist(rocObj@y.values)),4)

# adjust margins and plot TPR and FPR
par(mar = c(5,5,2,5))
plot(alpha,tpr, xlab="Threshold", xlim=c(0,1), ylab="True positive rate", type="l")
par(new="True")
plot(alpha,fpr, xlab="", ylab="", axes=F, xlim=c(0,1), type="l" )


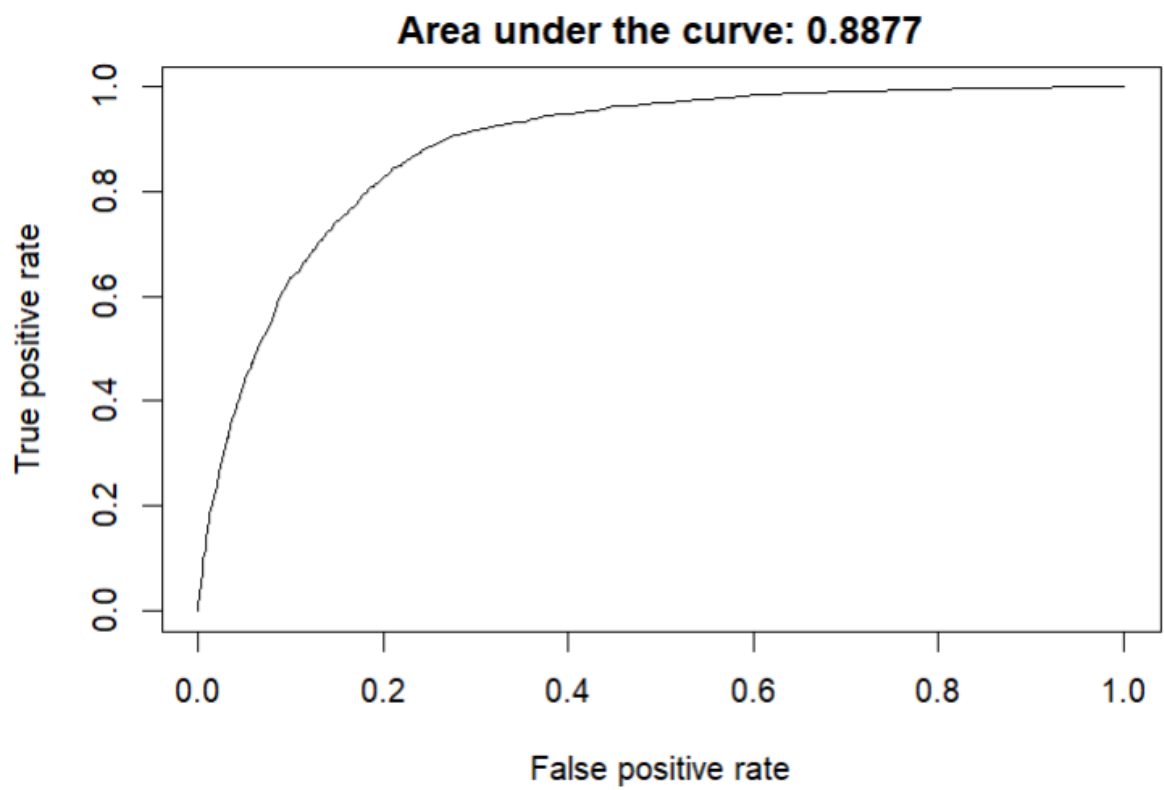
```

```
axis(side=4)
mtext(side=4, line=3, "False positive rate")
text(0.18,0.18,"FPR")
text(0.58,0.58,"TPR")
```

```
i <- which(round(alpha,2) == .5)
paste("Threshold=" , (alpha[i]) , " TPR=" , tpr[i] , " FPR=" , fpr[i])
```

```
i <- which(round(alpha,2) == .15)
paste("Threshold=" , (alpha[i]) , " TPR=" , tpr[i] , " FPR=" , fpr[i])
```

OUTPUT:



DECISION TREE

```
install.packages("rpart.plot") #install package rpart.plot
```

```
library("rpart")           #load libraries
```

```
library("rpart.plot")
```

```
play_decision <-
```

```
read.table("D:/DSA_DATASETS/datasets/banksample.csv",header=TRUE,sep=",")
```

```
play_decision
```

```
summary(play_decision)
```

```
x<- sort(runif(1000))
```

```
y<-data.frame(x=x,y=-x*log2(x)-(1-x)*log2(1-x))
```

```
plot(y,type="l",xlab="P(X=1)",ylab=expression("H"["X"]))
```

```
grid()
```

```
fit<rpарт(subscribed~job+marital+education+default+housing+loan+contact+poutcome,
```

```
      method="class",
```

```
      data=play_decision,
```

```
      control=rpart.control(minsplit=1),
```

```
      parms=list(split='information'))
```

```
summary(fit)
```

```
rpart.plot(fit, type=4,extra=2,clip.right.labs=FALSE,varlen=0,faclen=3)
```

```
newdata<-data.frame(job="retired",
```



```
marital="married",  
education="secondary",  
default="no",  
housing="yes",  
loan="no",  
contact="cellular",  
duration=598,  
poutcome="unknown"
```

```
)
```

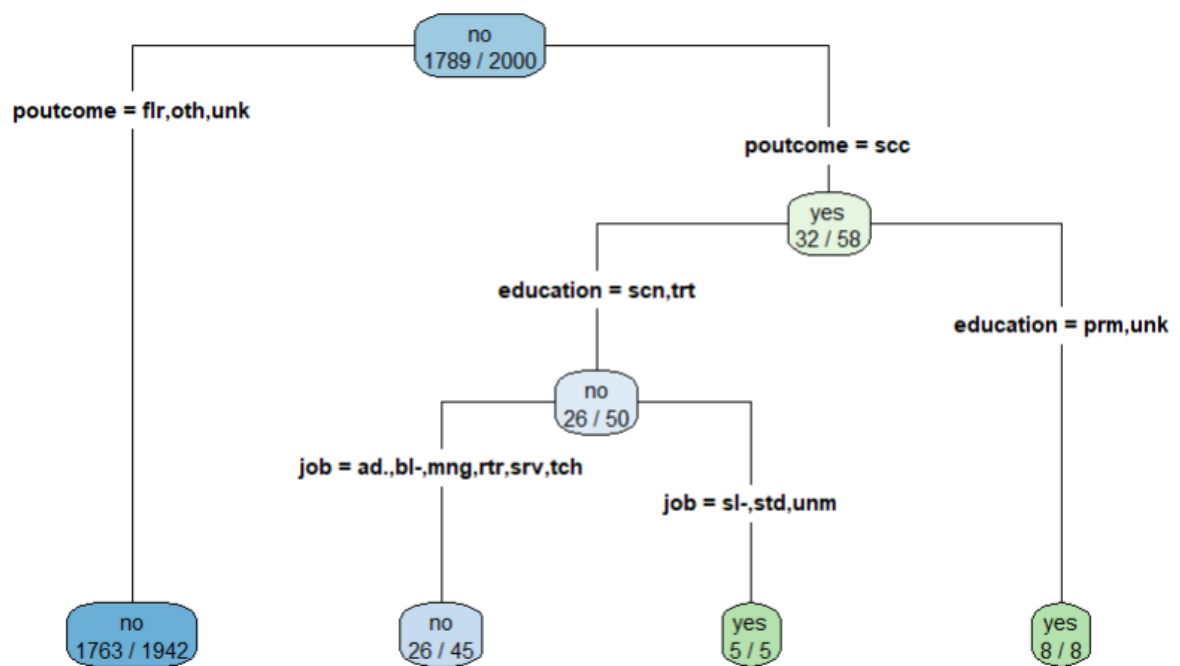

```
Newdata
```

```
predict(fit,newdata=newdata,type=c("class"))
```

```
library("rpart")      #load libraries
```

```
library("rpart.plot")
```

OUTPUT:



NAIVE BAYES CLASSIFIER

```
#####
```

```
# Naïve Bayes Classifier in R
```

```
#####
```

```
# Install required packages if not already installed
```

```
if(!require("e1071")) install.packages("e1071", dependencies=TRUE)
```

```
if(!require("rpart")) install.packages("rpart", dependencies=TRUE)
```

```
if(!require("rpart.plot")) install.packages("rpart.plot", dependencies=TRUE)
```

```
# Load necessary libraries
```

```
library(e1071)
```

```
library(rpart)
```

```
library(rpart.plot)
```

```
# Load dataset
```

```
banktrain <- read.csv("D:/DSA_DATASETS/datasets/bank-sample.csv", header=TRUE,  
sep=",")
```

```
# Drop unnecessary columns
```

```
drops <- c("balance", "day", "campaign", "pdays", "previous", "month")
```

```
banktrain <- banktrain[, !(names(banktrain) %in% drops)]
```

```
summary(banktrain)
```

```
# Train Naïve Bayes model
```

```
model <- naiveBayes(subscribed ~ ., data=banktrain)
```

```
print(model)
```

```
# Predict on training data
predictions <- predict(model, banktrain)
table(predictions, banktrain$subscribed)

# Load test data (uncomment if needed)
# banktest <- read.csv("D:/DSA_DATASETS/datasets/bank-sample-test.csv", header=TRUE,
# sep=",")
# banktest <- banktest[, !(names(banktest) %in% drops)]

# Predict on test data (if available)
# test_predictions <- predict(model, banktest)
# table(test_predictions, banktest$subscribed)

# Load another dataset for Naïve Bayes example
sample_data <- read.csv("D:/DSA_DATASETS/datasets/sample1.csv", header=TRUE,
sep=",")

# Split data into training and test sets
traindata <- sample_data[1:14, ]
testdata <- sample_data[15, ]

# Train Naïve Bayes model on sample dataset
model_sample <- naiveBayes(Enrolls ~ ., data=traindata)
print(model_sample)

# Predict on test sample
test_prediction <- predict(model_sample, testdata)
print(test_prediction)

# Train Naïve Bayes model with Laplace smoothing
```

```
model_sample_laplace <- naiveBayes(Enrolls ~ ., data=traindata, laplace=1)
print(model_sample_laplace)
```

```
# Predict using Laplace smoothing
```

```
laplace_prediction <- predict(model_sample_laplace, testdata)
print(laplace_prediction)
```

OUTPUT:

Naive Bayes Classifier for Discrete Predictors

Call:

```
naiveBayes.default(x = X, y = Y, laplace = laplace)
```

A-priori probabilities:

Y	No	Yes
	0.3571429	0.6428571

Conditional probabilities:

	Age		
Y	<=30	>40	31 to 40
No	0.8000000	0.6000000	0.2000000
Yes	0.3333333	0.4444444	0.5555556

	Income		
Y	High	Low	Medium
No	0.6000000	0.4000000	0.6000000
Yes	0.3333333	0.4444444	0.5555556

	JobSatisfaction	
Y	No	Yes
No	1.0000000	0.4000000
Yes	0.4444444	0.7777778

	Desire	
Y	Excellent	Fair
No	0.8000000	0.6000000
Yes	0.4444444	0.7777778

[1] Yes

Levels: No Yes

LINEAR REGRESSION

```
#####
```

```
# This code covers the code presented in
```

```
# Section 6.1 Linear Regression
```

```
#####
```

```
#####
```

```
# Section 6.1.2
```

```
#####
```

```
# Example in R
```

```
income_input=as.data.frame(read.csv("C:/Users/admin/Desktop/datasets/income.csv"))
```

```
income_input[1:10,]
```

```
summary(income_input)
```

```
library(lattice)
```

```
splom(~ income_input[c(2:5)], groups=NULL, data=income_input,
```

```
axis.line.tck = 0,
```

```
axis.text.alpha = 0)
```

```
results <- lm(Income ~ Age + Education + Gender, income_input)
```

```
summary(results)
```

```
results2 <- lm(Income ~ Age + Education, income_input)
```

```
summary(results2)
```

```
#####
```

```
# this code from the text is for illustrative purposes only
```

```
# the income_input variable does not contain the U.S. states
```

```
results3 <- lm(Income ~ Age + Education,
```

```
  + Alabama,
```

```
  + Alaska,
```

```
  + Arizona,
```

```
  + WestVirginia,
```

```
  + Wisconsin,
```

```
  income_input)
```

```
#####
```

```
# compute confidence intervals for the model parameters
```

```
confint(results2, level = .95)
```

```
# compute a confidence interval on the expected income of a person
```

```
Age <- 41
```

```
Education <- 12
```

```
new_pt <- data.frame(Age, Education)
```

```
conf_int_pt <- predict(results2, new_pt, level=.95, interval="confidence")
```

```
conf_int_pt
```

```
# compute a prediction interval on the income of the same person
```

```
pred_int_pt <- predict(results2, new_pt, level=.95, interval="prediction")
```

```
pred_int_pt
```




```
#####
```

```
# section 6.1.3 Diagnostics
```

```
#####
```

```
with(results2, {  
  plot(fitted.values, residuals,ylim=c(-40,40) )  
  points(c(min(fitted.values),max(fitted.values)), c(0,0), type = "l"))  
  
hist(results2$residuals, main="")  
  
qqnorm(results2$residuals, ylab="Residuals", main="")  
qqline(results2$residuals)
```

Outputs:



Scatter Plot Matrix

